




CPU SCHEDULING WINDOWS VS LINUX

Krish Sawant, Zohaib Quraishi,
Lyles Williams



What is CPU Scheduling?

- Process used in OS to decide which process gets to use CPU next
- Minimize response and waiting time of a process
- Preemptive Scheduling
 - a. OS can interrupt current process to have CPU handle another task
- Non-Preemptive Scheduling
 - a. Once process starts it has to finish before CPU handles another task
- Maximize CPU usage
- Different CPU Scheduling algorithms
 - a. RR(Round Robin)
 - b. FIFO(First come first serve)
 - c. SJF(Shortest Job First), etc
- Different OS like Windows and Linux have different algorithms for CPU scheduling



Why does CPU scheduling matter

- Affects how fast programs run and respond
- Ensures you are using your CPU efficiently
- Multiple processes which run at the same time have a fair and timely execution
- Ensures your hardware is being used to its fullest potential
- Makes sure hardware resources are not being wasted and are always doing something
- Without CPU Scheduling we could see things like
 - **Unoptimized servers**
 - **Inefficient CPU utilization**
 - **Hanging**
 - **A slow/unresponsive OS or program**
 - **A very high power consumption**
 - i. **Leads to high costs which could hurt a business**
- Windows and Linux each have their own strengths in their schedulers

Real World Use Cases


Windows Strengths

- Ensures for a smooth gaming experience with minimal lag and/or frame drops.
- When watching Netflix, Hulu, etc we will see not much lag or a slow, unresponsive website UI because of other less important tasks
- When doing a lot of work where one of many tabs are needed, the user will not have to worry about the computer slowing down.
- Windows tries to focus on their user experience along with responsiveness of the system



Linux Strengths

- When programming, we want to make sure all processes have equal resources while still using CPU efficiently.
- When we are running virtual machines we make sure each containers uses same amount of resources
- For some systems where all processes should have equal resources we should use CFS
- Linux tries to focus more on things like programming, servers, etc.



Windows Scheduler

Priority-Based Round Robin

- Combines preemptive Round Robin scheduling with Priorities
- Use a Priority queue in order to track priorities and which process will be executed next
 - a. **Priority of 1 means most important, the bigger the number, the less important**
- By combining preemptive Round Robin scheduling with Priorities we can achieve less average waiting time, average turnaround time and number of context switches
- Processes kept in increasing order of their remaining CPU burst time in the ready queue.
- New priorities are assigned according to the remaining CPU bursts of processes
- Process with shortest remaining CPU burst is assigned with highest priority

Round-Robin

- i. We go through all processes giving each a time quantum and keep cycling

Time quantum is 5ms

Process Name	CPU burst time (ms)	Priority
A	22	4
B	18	2
C	9	1
D	10	3
E	4	5

Burst time and priority for different processes

Table 2. Executed CPU burst for first round


S.No	Process	Executed Burst	Priority
1	C	5	1
2	B	5	2
3	D	5	3
4	A	5	4
5	E	4	5

Priority associated with each process


Table 3. Remaining CPU burst for second round & new assigned priorities

S.No	Process	Remaining Burst	Priority
1	C	4	1
2	D	5	2


3	B	13	3
4	A	17	4




Gantt chart for proposed algorithm



Context Switches



Average Wait Time





Average Turnaround Time


Linux Scheduler - CFS

- Default Linux Scheduler - Completely Fair Scheduler (CFS)
- Utilizes Red-black tree data structure
 - This keeps track of processes and CPU time
- Assigns CPU time to processes based on priority and amount of CPU time already used
- A general use scheduler
- Based on virtual runtime or ideal multitasking CPU
 - Virtual runtime is the amount of time a process has been running on the CPU
- Goal is to distribute all physical power and run each task at equal speeds in parallel
- Virtual runtime determines which process runs next.
 - Smallest virtual runtime runs next
 - Figured out through the red-black tree
- Nice values (from -20 to 19) determines which process gets priority to make things fair
 - -20 is the highest priority, 19 is the lowest
- Time Complexity: $O(\log n)$

CFS Tree Example





2 has the lowest virtual runtime so run it for a burst and remake the tree (bursts are 9 for this example)




5 has the lowest virtual runtime so run it for a burst and remake the tree (bursts are 9 for this example)

CFS Tree Example cont.



8 is the process with the lowest virtual runtime so run it for a burst (9)



This process continues until the processes terminate

Windows Scheduling vs Linux Scheduling Overview

	<u>Windows</u>	<u>Linux</u>
Scheduler Type	Priority Based Round Robin	Completely Fair Scheduler
Data Structure Used	Priority Queue	Red-black tree
Priority or Fairness	Priority Based	Fairness Based
Optimization	Optimized for responsiveness	General Use

Source Code Overview

- Simulated an environment where Priority Based Round Robin and Completely Fair Scheduler needed to run different processes of different priorities
- Kept track of variables like arrival, start, completion, waiting, response, average turn time, average wait time and average response time
- Made a Gantt chart

```
// -----
// Dummy workload
// -----
static Proc work[] = {
    // pid, arrival, burst, base_prio(0..15), nice(-20..19)
    {1, 0, 16, 10, 0}, // Medium priority → runs early but then yields
    {2, 2, 4, 8, -5}, // Moderate priority → runs after higher ones
    {3, 4, 20, 6, 5}, // Lowest priority → runs last
    {4, 6, 3, 12, -10}, // High priority → short interactive job → runs early
    {5, 10, 12, 7, 0}, // Lower priority → runs later
    {6, 12, 8, 14, 2}, // Very high priority → runs early
};
```

Simulation Results - Windows

```
===== Windows-like (Priority RR) =====
```

pid	arrival	burst	start	completion	waiting	response	Priority
1	0	16	1	31	15	1	10
2	2	4	32	36	30	30	8
3	4	20	51	73	49	47	6
4	6	3	22	25	16	16	12
5	10	12	37	50	28	27	7
6	12	8	13	21	1	1	14

Makespan=73 CPU_util=0.863 AvgTurn=33.67 AvgWait=23.17 AvgResp=20.33

```
==== Gantt: Windows-like ===
```

Time	P1	P6	P4	P1	P2	P5	P5	P3	P3	P3
0	Start									
12	End									
13		Start								
21		End								
22			Start							
25			End							
26				Start						
31				End						
32					Start					
36					End					
37						Start				
46						End				
47							Start			
50							End			
51								Start		
60								End		
61									Start	
70									End	
71										Start
73										End

- Low priority tasks wait longer
- Even if a high priority task arrives late it completes faster than low priority times
- Risk of starvation with the priority based round robin scheduling

Simulation Results - Linux

===== Linux CFS-like =====

pid	arrival	burst	start	completion	waiting	response	Priority
-----	---------	-------	-------	------------	---------	----------	----------

1	0	16	1	17	1	1	10
2	2	4	18	22	1	16	8
3	4	20	34	73	36	30	6
4	6	3	38	41	21	32	12
5	10	12	28	61	32	18	7
6	12	8	23	53	28	11	14

Makespan=73 CPU_util=0.863 AvgTurn=38.83 AvgWait=19.83 AvgResp=18.00

==== Gantt: Linux CFS-like ===

The Gantt chart displays the execution timeline for six processes (P1 through P6). The x-axis represents time, and the y-axis lists the processes. Each process's execution is shown as a horizontal bar. P1 starts at time 0 and ends at 17. P2 starts at 18 and ends at 22. P6 starts at 23 and ends at 27. P5 starts at 28 and ends at 33. P3 starts at 34 and ends at 37. P4 starts at 38 and ends at 41. P3 starts again at 42 and ends at 48. P6 starts again at 49 and ends at 53. P5 starts again at 54 and ends at 61. P3 starts again at 62 and ends at 73.

- CPU time distributed more evenly
- Waiting times are more fair
- High priority tasks are not prioritized

Simulation Results - Comparison & Conclusion

Comparison

- Linux CPU scheduling is more “fair”
- Windows works better if there are high priority tasks that need to be completed quickly

Conclusion

- Linux scheduling is more for general purpose while windows scheduling is more specialized

Work Cited

<https://documentation.ubuntu.com/real-time/latest/explanation/schedulers/>

https://www.geeksforgeeks.org/operating-systems/cpu-scheduling-in-operatin_g-systems/

<https://chiraggoyaliiit.medium.com/priority-based-round-robin-cpu-scheduling-algorithm-f9aa8517a844>